

✓ SQL at Scale with Spark SQL and DataFrames

Spark SQL brings native support for SQL to Spark and streamlines the process of querying data stored both in RDDs (Spark's distributed datasets) and in external sources. Spark SQL conveniently blurs the lines between RDDs and relational tables. Unifying these powerful abstractions makes it easy for developers to intermix SQL commands querying external data with complex analytics, all within a single application. Concretely, Spark SQL will allow developers to:

- Import relational data from Parquet files and Hive tables
- Run SQL queries over imported data and existing RDDs
- Easily write RDDs out to Hive tables or Parquet files

Spark SQL also includes a cost-based optimizer, columnar storage, and code generation to make queries fast. At the same time, it scales to thousands of nodes and multi-hour queries using the Spark engine, which provides full mid-query fault tolerance, without having to worry about using a different engine for historical data.

For getting a deeper perspective into the background, concepts, architecture of Spark SQL and DataFrames you can check out the original article, ['_SQL at Scale with Apache Spark SQL and DataFrames-Concepts, Architecture and Examples'_](#)

This tutorial will familiarize you with essential Spark capabilities to deal with structured data typically often obtained from databases or flat files. We will explore typical ways of querying and aggregating relational data by leveraging concepts of DataFrames and SQL using Spark. We will work on an interesting dataset from the [KDD Cup 1999](#) and try to query the data using high level abstractions like the dataframe which has already been a hit in popular data analysis tools like R and Python. We will also look at how easy it is to build data queries using the SQL language which you have learnt and retrieve insightful information from our data. This also happens at scale without us having to do a lot more since Spark distributes these data structures efficiently in the backend which makes our queries scalable and as efficient as possible.

```
1 import pandas as pd
2 import matplotlib.pyplot as plt
3 plt.style.use('fivethirtyeight')

1 import pyspark
2 from pyspark.sql import SparkSession
3
4 spark = SparkSession.builder.appName("nik").getOrCreate()
5
6 from pyspark import SparkContext
7 sc = SparkContext.getOrCreate()
```

Data Retrieval

We will use data from the [KDD Cup 1999](#), which is the data set used for The Third International Knowledge Discovery and Data Mining Tools Competition, which was held in conjunction with KDD-99 The Fifth International Conference on Knowledge Discovery and Data Mining. The competition task was to build a network intrusion detector, a predictive model capable of distinguishing between "bad" connections, called intrusions or attacks, and "good" normal connections. This database contains a standard set of data to be audited, which includes a wide variety of intrusions simulated in a military network environment.

We will be using the reduced dataset `kddcup.data_10_percent.gz` containing nearly half million network interactions since we would be downloading this Gzip file from the web locally and then work on the same. If you have a good, stable internet connection, feel free to download and work with the full dataset available as `kddcup.data.gz`

✓ Building the KDD Dataset

Now that we have our data stored in the Databricks filesystem. Let's load up our data from the disk into Spark's traditional abstracted data structure, the RDD

```
1 data_file = "kddcup_data.gz"
2 raw_rdd = sc.textFile(data_file)
```

```
1 print(raw_rdd)
```

```
↗ kddcup_data.gz MapPartitionsRDD[1] at textFile at NativeMethodAccessorImpl.java:0
```

```
1 raw_rdd.collect()
```



```
'0,tcp,http,SF,246,1718,0,0,0,0,0,1,0,0,0,0,0,0,0,0,10,12,0.00,0.00,0.00,0.00,1.00,0.00,0.25,255,255,1.00,0.00,0.00,0.00,0.00
'0,tcp,http,SF,218,1484,0,0,0,0,0,1,0,0,0,0,0,0,0,0,0,7,7,0.00,0.00,0.00,0.00,1.00,0.00,0.00,255,255,1.00,0.00,0.00,0.00,0.00,0
'0,tcp,http,SF,222,1651,0,0,0,0,0,1,0,0,0,0,0,0,0,0,0,4,4,0.00,0.00,0.00,0.00,1.00,0.00,0.00,255,255,1.00,0.00,0.00,0.00,0.00,0
'0,tcp,http,SF,215,1108,0,0,0,0,0,1,0,0,0,0,0,0,0,0,0,1,1,0.00,0.00,0.00,0.00,1.00,0.00,0.00,255,255,1.00,0.00,0.00,0.00,0.00,0
'0,tcp,http,SF,216,1645,0,0,0,0,0,1,0,0,0,0,0,0,0,0,0,3,8,0.00,0.00,0.00,0.00,1.00,0.00,0.50,255,255,1.00,0.00,0.00,0.00,0.00,0
'0,tcp,http,SF,213,1694,0,0,0,0,0,1,0,0,0,0,0,0,0,0,0,13,18,0.00,0.00,0.00,0.00,1.00,0.00,0.22,255,255,1.00,0.00,0.00,0.00,0.00,0
'0,tcp,http,SF,243,1719,0,0,0,0,0,1,0,0,0,0,0,0,0,0,0,6,24,0.00,0.00,0.00,0.00,1.00,0.00,0.08,255,255,1.00,0.00,0.00,0.00,0.00,0
'0,tcp,http,SF,147,3253,0,0,0,0,0,1,0,0,0,0,0,0,0,0,0,3,3,0.00,0.00,0.00,0.00,1.00,0.00,0.00,255,255,1.00,0.00,0.00,0.00,0.00,0
'0,tcp,http,SF,182,4023,0,0,0,0,0,1,0,0,0,0,0,0,0,0,0,9,9,0.00,0.00,0.00,0.00,1.00,0.00,0.00,255,255,1.00,0.00,0.00,0.00,0.00,0
'0,tcp,http,SF,195,24572,0,0,0,0,0,1,0,0,0,0,0,0,0,0,0,2,2,0.00,0.00,0.00,0.00,1.00,0.00,0.00,255,255,1.00,0.00,0.00,0.00,0.00,0
'0,tcp,smtp,SF,3366,329,0,0,0,0,0,1,0,0,0,0,0,0,0,0,0,1,1,0.00,0.00,0.00,0.00,1.00,0.00,0.00,8,7,0.88,0.25,0.12,0.00,0.00,0.00,
'0,tcp,smtp,SF,1666,328,0,0,0,0,0,1,0,0,0,0,0,0,0,0,0,1,1,0.00,0.00,0.00,0.00,1.00,0.00,0.00,18,12,0.67,0.17,0.06,0.00,0.00,0.0
'0,tcp,smtp,SF,751,279,0,0,0,0,0,1,0,0,0,0,0,0,0,0,0,1,1,0.00,0.00,0.00,0.00,1.00,0.00,0.00,28,19,0.68,0.14,0.04,0.00,0.00,0.00
'0,tcp,finger,SF,9,140,0,0,0,0,0,0,0,0,0,0,0,0,0,0,0,1,1,0.00,0.00,0.00,0.00,1.00,0.00,0.00,38,5,0.13,0.11,0.03,0.00,0.00,0.00,
'0,tcp,smtp,SF,620,329,0,0,0,0,0,1,0,0,0,0,0,0,0,0,0,1,2,0.00,0.00,0.00,0.00,1.00,0.00,1.00,48,26,0.54,0.10,0.02,0.00,0.00,0.00
'0,tcp,smtp,SF,805,327,0,0,0,0,0,1,0,0,0,0,0,0,0,0,0,2,2,0.00,0.00,0.00,0.00,1.00,0.00,0.00,58,35,0.60,0.09,0.02,0.00,0.00,0.00
'0,tcp,smtp,SF,842,279,0,0,0,0,0,1,0,0,0,0,0,0,0,0,0,2,1,0.00,0.00,0.00,0.00,0.50,1.00,0.00,68,41,0.60,0.07,0.01,0.00,0.00,0.00
'0,udp,domain_u,SF,33,0,0,0,0,0,0,0,0,0,0,0,0,0,0,0,0,1,2,0.00,0.00,0.00,0.00,1.00,0.00,1.00,78,14,0.18,0.06,0.18,0.00,0.00,0.0
'0,udp,domain_u,SF,30,0,0,0,0,0,0,0,0,0,0,0,0,0,0,0,0,1,1,0.00,0.00,0.00,0.00,1.00,0.00,0.00,88,15,0.17,0.06,0.17,0.00,0.00,0.0
'0,tcp,smtp,SF,1375,328,0,0,0,0,0,1,0,0,0,0,0,0,0,0,0,1,1,0.00,0.00,0.00,0.00,1.00,0.00,0.00,98,64,0.65,0.05,0.01,0.00,0.00,0.0
'1,tcp,smtp,SF,1550,326,0,0,0,0,0,1,0,0,0,0,0,0,0,0,0,1,3,0.00,0.00,0.00,0.00,1.00,0.00,1.00,108,72,0.67,0.05,0.01,0.00,0.00,0.
'0,udp,domain_u,SF,30,0,0,0,0,0,0,0,0,0,0,0,0,0,0,0,0,1,4,0.00,0.00,0.00,0.00,1.00,0.00,0.75,118,21,0.18,0.04,0.18,0.00,0.00,0.
'0,tcp,smtp,SF,1254,323,0,0,0,0,0,1,0,0,0,0,0,0,0,0,0,1,1,0.00,0.00,0.00,0.00,1.00,0.00,0.00,5,82,0.80,0.40,0.20,0.02,0.00,0.00
'0,udp,domain_u,SF,31,0,0,0,0,0,0,0,0,0,0,0,0,0,0,0,0,1,2,0.00,0.00,0.00,0.00,1.00,0.00,1.00,15,26,0.27,0.20,0.27,0.08,0.00,0.0
'0,tcp,smtp,SF,2753,280,0,0,0,0,0,1,0,0,0,0,0,0,0,0,0,1,1,0.00,0.00,0.00,0.00,1.00,0.00,0.00,25,95,0.68,0.16,0.04,0.02,0.00,0.0
'0,tcp,auth,SF,9,35,0,0,0,0,0,1,0,0,0,0,0,0,0,0,0,3,1,0.00,0.00,0.00,0.00,0.33,0.67,0.00,35,12,0.11,0.11,0.03,0.17,0.00,0.00,0.
```

```
1 raw_rdd.count()
```

```
494021
```

```
1 type(raw_rdd)
```

```
pyspark.rdd.RDD
```

Building a Spark DataFrame on our Data

A Spark DataFrame is an interesting data structure representing a distributed collection of data. A DataFrame is a Dataset organized into named columns. It is conceptually equivalent to a table in a relational database or a dataframe in R/Python, but with richer optimizations under the hood. DataFrames can be constructed from a wide array of sources such as: structured data files, tables in Hive, external databases, or existing RDDs in our case.

Typically the entry point into all SQL functionality in Spark is the `SQLContext` class. To create a basic instance of this call, all we need is a `SparkContext` reference. In Databricks this global context object is available as `sc` for this purpose.

```
1 from pyspark.sql import SQLContext
2 sqlContext = SQLContext(sc)
3 sqlContext
```

```
C:\spark\python\pyspark\sql\context.py:113: FutureWarning: Deprecated in 3.0.0. Use SparkSession.builder.getOrCreate() instead.
  warnings.warn(
<pyspark.sql.context.SQLContext at 0x1c50a45a810>
```

- Splitting the CSV data

Each entry in our RDD is a comma-separated line of data which we first need to split before we can parse and build our dataframe

```
1 csv_rdd = raw_rdd.map(lambda row: row.split(","))
2
```

```
1 csv_rdd.collect()
```

```
1 print(csv_rdd.take(2))
2
```

[[['0', 'tcp', 'http', 'SF', '181', '5450', '0', '0', '0', '0', '0', '1', '0', '0', '0', '0', '0', '0', '0', '0', '0', '0', '0', '0', '8', '8',

```
1 print(type(csv_rdd))
2
```

```
>>> rdd
Out[10]: <class 'pyspark.rdd.PipelinedRDD'>
```

- ✓ Check the total number of features (columns)

We can use the following code to check the total number of potential columns in our dataset

```
1 len(csv_rdd.take(1)[0])
```

 42

▼ Data Understanding and Parsing

The KDD 99 Cup data consists of different attributes captured from connection data. The full list of attributes in the data can be obtained [here](#) and further details pertaining to the description for each attribute\column can be found [here](#). We will just be using some specific columns from the dataset, the details of which are specified below.

feature num	feature name	description	type
1	duration	length (number of seconds) of the connection	continuous
2	protocol_type	type of the protocol, e.g. tcp, udp, etc.	discrete
3	service	network service on the destination, e.g., http, telnet, etc.	discrete
4	src_bytes	number of data bytes from source to destination	continuous
5	dst_bytes	number of data bytes from destination to source	continuous
6	flag	normal or error status of the connection	discrete
7	wrong_fragment	number of "wrong" fragments	continuous
8	urgent	number of urgent packets	continuous
9	hot	number of "hot" indicators	continuous
10	num_failed_logins	number of failed login attempts	continuous
11	num_compromised	number of "compromised" conditions	continuous
12	su_attempted	1 if "su root" command attempted; 0 otherwise	discrete
13	num_root	number of "root" accesses	continuous
14	num_file_creations	number of file creation operations	continuous

We will be extracting the following columns based on their positions in each datapoint (row) and build a new RDD as follows

```
1 from pyspark.sql import Row
2
3 parsed_rdd = csv_rdd.map(lambda r: Row(
4     duration=int(r[0]),
5     protocol_type=r[1],
6     service=r[2],
7     flag=r[3],
8     src_bytes=int(r[4]),
9     dst_bytes=int(r[5]),
10    wrong_fragment=int(r[7]),
11    urgent=int(r[8]),
12    hot=int(r[9]),
13    num_failed_logins=int(r[10]),
14    num_compromised=int(r[12]),
15    su_attempted=r[14],
16    num_root=int(r[15]),
17    num_file_creations=int(r[16]),
18    label=r[-1]
19 )
20 )
21
```

```
1 parsed_rdd.collect()
```



```

Row(duration=0, protocol_type= tcp , service= nntp , flag= SF , src_bytes=305, dst_bytes=19509, wrong_fragment=0, urgent=0,
hot=0, num_failed_logins=0, num_compromised=0, su_attempted='0', num_root=0, num_file_creations=0, label='normal.'),
Row(duration=0, protocol_type='tcp', service='http', flag='SF', src_bytes=231, dst_bytes=1869, wrong_fragment=0, urgent=0,
hot=0, num_failed_logins=0, num_compromised=0, su_attempted='0', num_root=0, num_file_creations=0, label='normal.'),
Row(duration=0, protocol_type='tcp', service='http', flag='SF', src_bytes=217, dst_bytes=363, wrong_fragment=0, urgent=0,
hot=0, num_failed_logins=0, num_compromised=0, su_attempted='0', num_root=0, num_file_creations=0, label='normal.'),
Row(duration=0, protocol_type='tcp', service='http', flag='SF', src_bytes=297, dst_bytes=2769, wrong_fragment=0, urgent=0,
hot=0, num_failed_logins=0, num_compromised=0, su_attempted='0', num_root=0, num_file_creations=0, label='normal.'),
Row(duration=0, protocol_type='tcp', service='http', flag='SF', src_bytes=308, dst_bytes=1278, wrong_fragment=0, urgent=0,
hot=0, num_failed_logins=0, num_compromised=0, su_attempted='0', num_root=0, num_file_creations=0, label='normal.'),
Row(duration=0, protocol_type='tcp', service='http', flag='SF', src_bytes=274, dst_bytes=196, wrong_fragment=0, urgent=0,
hot=0, num_failed_logins=0, num_compromised=0, su_attempted='0', num_root=0, num_file_creations=0, label='normal.'),
Row(duration=0, protocol_type='tcp', service='http', flag='SF', src_bytes=291, dst_bytes=5657, wrong_fragment=0, urgent=0,
hot=0, num_failed_logins=0, num_compromised=0, su_attempted='0', num_root=0, num_file_creations=0, label='normal.'),
Row(duration=0, protocol_type='tcp', service='http', flag='SF', src_bytes=328, dst_bytes=1433, wrong_fragment=0, urgent=0,
hot=0, num_failed_logins=0, num_compromised=0, su_attempted='0', num_root=0, num_file_creations=0, label='normal.'),
Row(duration=0, protocol_type='tcp', service='http', flag='SF', src_bytes=211, dst_bytes=2101, wrong_fragment=0, urgent=0,
hot=0, num_failed_logins=0, num_compromised=0, su_attempted='0', num_root=0, num_file_creations=0, label='normal.'),
Row(duration=0, protocol_type='tcp', service='http', flag='SF', src_bytes=325, dst_bytes=2452, wrong_fragment=0, urgent=0,
hot=0, num_failed_logins=0, num_compromised=0, su_attempted='0', num_root=0, num_file_creations=0, label='normal.'),
Row(duration=0, protocol_type='tcp', service='http', flag='SF', src_bytes=340, dst_bytes=22965, wrong_fragment=0, urgent=0,
hot=0, num_failed_logins=0, num_compromised=0, su_attempted='0', num_root=0, num_file_creations=0, label='normal.'),
Row(duration=0, protocol_type='tcp', service='http', flag='SF', src_bytes=235, dst_bytes=1028, wrong_fragment=0, urgent=0,
hot=0, num_failed_logins=0, num_compromised=0, su_attempted='0', num_root=0, num_file_creations=0, label='normal.'),
Row(duration=0, protocol_type='tcp', service='http', flag='SF', src_bytes=223, dst_bytes=1275, wrong_fragment=0, urgent=0,
hot=0, num_failed_logins=0, num_compromised=0, su_attempted='0', num_root=0, num_file_creations=0, label='normal.'),
Row(duration=0, protocol_type='tcp', service='http', flag='SF', src_bytes=226, dst_bytes=4006, wrong_fragment=0, urgent=0,
hot=0, num_failed_logins=0, num_compromised=0, su_attempted='0', num_root=0, num_file_creations=0, label='normal.'),
Row(duration=0, protocol_type='tcp', service='http', flag='SF', src_bytes=231, dst_bytes=5283, wrong_fragment=0, urgent=0,
hot=1, num_failed_logins=0, num_compromised=0, su_attempted='0', num_root=0, num_file_creations=0, label='normal.'),
Row(duration=0, protocol_type='tcp', service='http', flag='SF', src_bytes=230, dst_bytes=5558, wrong_fragment=0, urgent=0,
hot=0, num_failed_logins=0, num_compromised=0, su_attempted='0', num_root=0, num_file_creations=0, label='normal.'),
1

```

```
1 parsed_rdd.take(5)
```



```

-----
Py4JJavaError                                Traceback (most recent call last)
Cell In[20], line 1
----> 1 parsed_rdd.take(5)

File C:\spark\python\pyspark\rdd.py:2855, in RDD.take(self, num)
    2852         taken += 1
    2854 p = range(partsScanned, min(partsScanned + numPartsToTry, totalParts))
-> 2855 res = self.context.runJob(self, takeUpToNumLeft, p)
    2857 items += res
    2858 partsScanned += numPartsToTry

File C:\spark\python\pyspark\context.py:2510, in SparkContext.runJob(self, rdd, partitionFunc, partitions, allowLocal)
    2508 mappedRDD = rdd.mapPartitions(partitionFunc)
    2509 assert self._jvm is not None
-> 2510 sock_info = self._jvm.PythonRDD.runJob(self._jsc.sc(), mappedRDD._jrdd, partitions)
    2511 return list(_load_from_socket(sock_info, mappedRDD._jrdd_deserializer))

File C:\python\Lib\site-packages\py4j\java_gateway.py:1322, in JavaMember.__call__(self, *args)
    1316 command = proto.CALL_COMMAND_NAME + \
    1317     self.command_header + \
    1318     args_command + \
    1319     proto.END_COMMAND_PART
    1321 answer = self.gateway_client.send_command(command)
-> 1322 return_value = get_return_value(
    1323     answer, self.gateway_client, self.target_id, self.name)
    1325 for temp_arg in temp_args:
    1326     if hasattr(temp_arg, "_detach"):

File C:\spark\python\pyspark\errors\exceptions\captured.py:179, in capture_sql_exception.<locals>.deco(*a, **kw)
    177 def deco(*a: Any, **kw: Any) -> Any:
    178     try:
--> 179         return f(*a, **kw)
    180     except Py4JJavaError as e:
    181         converted = convert_exception(e.java_exception)

File C:\python\Lib\site-packages\py4j\protocol.py:326, in get_return_value(answer, gateway_client, target_id, name)
    324 value = OUTPUT_CONVERTER[type](answer[2:], gateway_client)
    325 if answer[1] == REFERENCE_TYPE:
--> 326     raise Py4JJavaError(
    327         "An error occurred while calling {0}{1}{2}.\n".
    328         format(target_id, ".", name), value)
    329 else:
    330     raise Py4JError(
    331         "An error occurred while calling {0}{1}{2}. Trace:\n{3}\n".
    332         format(target_id, ".", name, value))

```

Py4JJavaError: An error occurred while calling z:org.apache.spark.api.python.PythonRDD.runJob.
: org.apache.spark.SparkException: Job aborted due to stage failure: Task 0 in stage 10.0 failed 1 times, most recent failure:
Lost task 0.0 in stage 10.0 (TID 10) (konoha executor driver): java.net.SocketException: Connection reset by peer: socket write
error

```

at java.net.SocketOutputStream.socketWrite0(Native Method)
at java.net.SocketOutputStream.socketWrite(SocketOutputStream.java:111)
at java.net.SocketOutputStream.write(SocketOutputStream.java:155)
at java.io.BufferedOutputStream.flushBuffer(BufferedOutputStream.java:82)
at java.io.BufferedOutputStream.write(BufferedOutputStream.java:126)
at java.io.DataOutputStream.write(DataOutputStream.java:107)
at java.io.FilterOutputStream.write(FilterOutputStream.java:97)
at org.apache.spark.api.python.PythonRDD$.writeUTF(PythonRDD.scala:492)
at org.apache.spark.api.python.PythonRDD$.write$1(PythonRDD.scala:312)
at org.apache.spark.api.python.PythonRDD$.anonfun$writeIteratorToStream$1(PythonRDD.scala:322)
at org.apache.spark.api.python.PythonRDD$.anonfun$writeIteratorToStream$1$adapted(PythonRDD.scala:322)
at scala.collection.IterableOnceOps.foreach(IterableOnce.scala:563)
at scala.collection.IterableOnceOps.foreach$(IterableOnce.scala:561)
at scala.collection.AbstractIterator.foreach(Iterator.scala:1293)
at org.apache.spark.api.python.PythonRDD$.writeIteratorToStream(PythonRDD.scala:322)
at org.apache.spark.api.python.PythonRunner$.anonfun$2.writeIteratorToStream(PythonRunner.scala:751)
at org.apache.spark.api.python.BasePythonRunner$WriterThread$.anonfun$run$1(PythonRunner.scala:451)
at org.apache.spark.util.Utils$.logUncaughtExceptions(Utils.scala:1928)
at org.apache.spark.api.python.BasePythonRunner$WriterThread.run(PythonRunner.scala:282)

```

Driver stacktrace:

```

at org.apache.spark.scheduler.DAGScheduler.failJobAndIndependentStages(DAGScheduler.scala:2844)
at org.apache.spark.scheduler.DAGScheduler$.anonfun$abortStage$2(DAGScheduler.scala:2780)
at org.apache.spark.scheduler.DAGScheduler$.anonfun$abortStage$2$adapted(DAGScheduler.scala:2779)
at scala.collection.immutable.List.foreach(List.scala:333)
at org.apache.spark.scheduler.DAGScheduler.abortStage(DAGScheduler.scala:2779)
at org.apache.spark.scheduler.DAGScheduler$.anonfun$handleTaskSetFailed$1(DAGScheduler.scala:1242)
at org.apache.spark.scheduler.DAGScheduler$.anonfun$handleTaskSetFailed$1$adapted(DAGScheduler.scala:1242)
at scala.Option.foreach(Option.scala:437)
at org.apache.spark.scheduler.DAGScheduler.handleTaskSetFailed(DAGScheduler.scala:1242)
at org.apache.spark.scheduler.DAGSchedulerEventProcessLoop.doOnReceive(DAGScheduler.scala:3048)
at org.apache.spark.scheduler.DAGSchedulerEventProcessLoop.onReceive(DAGScheduler.scala:2982)
at org.apache.spark.scheduler.DAGSchedulerEventProcessLoop.onReceive(DAGScheduler.scala:2971)
at org.apache.spark.util.EventLoop$.anonfun$1.run(EventLoop.scala:49)
at org.apache.spark.scheduler.DAGScheduler.runJob(DAGScheduler.scala:984)
at org.apache.spark.SparkContext.runJob(SparkContext.scala:2398)
at org.apache.spark.SparkContext.runJob(SparkContext.scala:2419)
at org.apache.spark.SparkContext.runJob(SparkContext.scala:2438)
at org.apache.spark.api.python.PythonRDD$.runJob(PythonRDD.scala:181)

```

```

at org.apache.spark.api.python.PythonRDD.runJob(PythonRDD.scala)
at sun.reflect.NativeMethodAccessorImpl.invoke0(Native Method)
at sun.reflect.NativeMethodAccessorImpl.invoke(NativeMethodAccessorImpl.java:62)
at sun.reflect.DelegatingMethodAccessorImpl.invoke(DelegatingMethodAccessorImpl.java:43)
at java.lang.reflect.Method.invoke(Method.java:498)
at py4j.reflection.MethodInvoker.invoke(MethodInvoker.java:244)
at py4j.reflection.ReflectionEngine.invoke(ReflectionEngine.java:374)
at py4j.Gateway.invoke(Gateway.java:282)
at py4j.commands.AbstractCommand.invokeMethod(AbstractCommand.java:132)
at py4j.commands.CallCommand.execute(CallCommand.java:79)
at py4j.ClientServerConnection.waitForCommands(ClientServerConnection.java:182)
at py4j.ClientServerConnection.run(ClientServerConnection.java:106)
at java.lang.Thread.run(Thread.java:748)
Caused by: java.net.SocketException: Connection reset by peer: socket write error
at java.net.SocketOutputStream.socketWrite0(Native Method)
at java.net.SocketOutputStream.socketWrite(SocketOutputStream.java:111)
at java.net.SocketOutputStream.write(SocketOutputStream.java:155)
at java.io.BufferedOutputStream.flushBuffer(BufferedOutputStream.java:82)
at java.io.BufferedOutputStream.write(BufferedOutputStream.java:126)
at java.io.DataOutputStream.write(DataOutputStream.java:107)
at java.io.FilterOutputStream.write(FilterOutputStream.java:97)
at org.apache.spark.api.python.PythonRDD$.writeUTF(PythonRDD.scala:492)
at org.apache.spark.api.python.PythonRDD$.write$1(PythonRDD.scala:312)
at org.apache.spark.api.python.PythonRDD$.anonfun$writeIteratorToStream$1(PythonRDD.scala:322)
at org.apache.spark.api.python.PythonRDD$.anonfun$writeIteratorToStream$1$adapted(PythonRDD.scala:322)
at scala.collection.IterableOnceOps.foreach(IterableOnce.scala:563)
at scala.collection.IterableOnceOps.foreach$(IterableOnce.scala:561)
at scala.collection.AbstractIterator.foreach(Iterator.scala:1293)
at org.apache.spark.api.python.PythonRDD$.writeIteratorToStream(PythonRDD.scala:322)
at org.apache.spark.api.python.PythonRunner$.anonfun$2.writeIteratorToStream(PythonRunner.scala:751)
at org.apache.spark.api.python.BasePythonRunner$WriterThread$.anonfun$run$1(PythonRunner.scala:451)
at org.apache.spark.util.Utils$.logUncaughtExceptions(Utils.scala:1928)
at org.apache.spark.api.python.BasePythonRunner$WriterThread.run(PythonRunner.scala:282)

```

✓ Constructing the DataFrame

Now that our data is neatly parsed and formatted, let's build our DataFrame!

```

1 df = sqlContext.createDataFrame(parsed_rdd)
2 df.show(10)

```

```

+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
|duration|protocol_type|service|flag|src_bytes|dst_bytes|wrong_fragment|urgent|hot|num_failed_logins|num_compromised|su_attempted|num
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
|0|0|tcp|http|SF|181|5450|0|0|0|0|0|0|
|0|0|tcp|http|SF|239|486|0|0|0|0|0|0|
|0|0|tcp|http|SF|235|1337|0|0|0|0|0|0|
|0|0|tcp|http|SF|219|1337|0|0|0|0|0|0|
|0|0|tcp|http|SF|217|2032|0|0|0|0|0|0|
|0|0|tcp|http|SF|217|2032|0|0|0|0|0|0|
|0|0|tcp|http|SF|212|1940|0|0|0|0|0|0|
|0|0|tcp|http|SF|159|4087|0|0|0|0|0|0|
|0|0|tcp|http|SF|210|151|0|0|0|0|0|0|
|0|0|tcp|http|SF|212|786|0|0|1|0|0|0|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
only showing top 10 rows

```

Now, we can easily have a look at our dataframe's schema using the `printSchema(...)` function.

```
1 df.printSchema()
```

```

root
|-- duration: long (nullable = true)
|-- protocol_type: string (nullable = true)
|-- service: string (nullable = true)
|-- flag: string (nullable = true)
|-- src_bytes: long (nullable = true)
|-- dst_bytes: long (nullable = true)
|-- wrong_fragment: long (nullable = true)
|-- urgent: long (nullable = true)
|-- hot: long (nullable = true)
|-- num_failed_logins: long (nullable = true)
|-- num_compromised: long (nullable = true)
|-- su_attempted: string (nullable = true)
|-- num_root: long (nullable = true)
|-- num_file_creations: long (nullable = true)
|-- label: string (nullable = true)

```

✓ Building a temporary table

We can leverage the `registerTempTable()` function to build a temporary table to run SQL commands on our DataFrame at scale! A point to remember is that the lifetime of this temp table is tied to the session. It creates an in-memory table that is scoped to the cluster in which it was created. The data is stored using Hive's highly-optimized, in-memory columnar format.

You can also check out `saveAsTable()` which creates a permanent, physical table stored in S3 using the Parquet format. This table is accessible to all clusters. The table metadata including the location of the file(s) is stored within the Hive metastore.

```
1 help(df.registerTempTable)
```

 Help on method registerTempTable in module pyspark.sql.dataframe:

```
registerTempTable(name: str) -> None method of pyspark.sql.dataframe.DataFrame instance
Registers this :class:`DataFrame` as a temporary table using the given name.
```

```
The lifetime of this temporary table is tied to the :class:`SparkSession`
that was used to create this :class:`DataFrame`.
```

```
.. versionadded:: 1.3.0
```

```
.. versionchanged:: 3.4.0
   Supports Spark Connect.
```

```
.. deprecated:: 2.0.0
   Use :meth:`DataFrame.createOrReplaceTempView` instead.
```

```
Parameters
```

```
-----
```

```
name : str
```

```
    Name of the temporary table to register.
```

```
Examples
```

```
-----
```

```
>>> df = spark.createDataFrame([(2, "Alice"), (5, "Bob")], schema=["age", "name"])
>>> df.registerTempTable("people")
>>> df2 = spark.sql("SELECT * FROM people")
>>> sorted(df.collect()) == sorted(df2.collect())
True
>>> spark.catalog.dropTempView("people")
True
```

```
1 df.registerTempTable("connections")
```

 C:\spark\python\pyspark\sql\dataframe.py:329: FutureWarning: Deprecated in 2.0, use createOrReplaceTempView instead. warnings.warn("Deprecated in 2.0, use createOrReplaceTempView instead.", FutureWarning)


✓ Executing SQL at Scale

Let's look at a few examples of how we can run SQL queries on our table based off our dataframe. We will start with some simple queries and then look at aggregations, filters, sorting, subqueries and pivots

✓ Connections based on the protocol type

Let's look at how we can get the total number of connections based on the type of connectivity protocol. First we will get this information using normal DataFrame DSL syntax to perform aggregations.

```
1 df.groupBy('protocol_type').count().orderBy('count', ascending=False).show()
```

 +-----+-----+
|protocol_type| count|
+-----+-----+
icmp	283602
tcp	190065
udp	20354
+-----+-----+

Can we also use SQL to perform the same aggregation? Yes we can leverage the table we built earlier for this!

```
1 protocols = sqlContext.sql("""
2     SELECT protocol_type, count(*) as freq
3     FROM connections
4     GROUP BY protocol_type
5     ORDER BY 2 DESC
6     """)
7 protocols.show()
```



```

+-----+-----+
|protocol_type| freq|
+-----+-----+
|      icmp|283602|
|      tcp|190065|
|      udp| 20354|
+-----+-----+

```

You can clearly see, that you get the same results and you do not need to worry about your background infrastructure or how the code is executed. Just write simple SQL!

✓ Connections based on good or bad (attack types) signatures

We will now run a simple aggregation to check the total number of connections based on good (normal) or bad (intrusion attacks) types.

```

1 labels = sqlContext.sql("""
2                               SELECT label, count(*) as freq
3                               FROM connections
4                               GROUP BY label
5                               ORDER BY 2 DESC
6                               """)
7 labels.show()

```

```

+-----+-----+
|      label| freq|
+-----+-----+
|      smurf.|280790|
|      neptune.|107201|
|      normal.| 97278|
|      back.| 2203|
|      satan.| 1589|
|      ipsweep.| 1247|
|      portsweep.| 1040|
|      warezclient.| 1020|
|      teardrop.| 979|
|      pod.| 264|
|      nmap.| 231|
|      guess_passwd.| 53|
|      buffer_overflow.| 30|
|      land.| 21|
|      warezmaster.| 20|
|      imap.| 12|
|      rootkit.| 10|
|      loadmodule.| 9|
|      ftp_write.| 8|
|      multihop.| 7|
+-----+-----+
only showing top 20 rows

```

```

1 labels = sqlContext.sql("""
2                               SELECT label, count(*) as freq
3                               FROM connections
4                               GROUP BY label
5                               ORDER BY 2 DESC
6                               """)
7 labels.show()

```

```

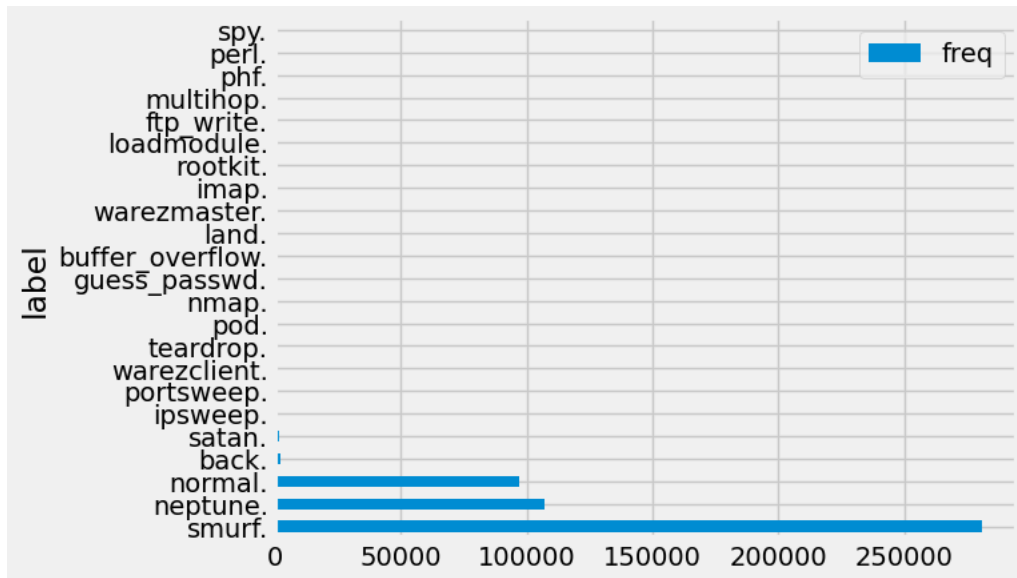
+-----+-----+
|      label| freq|
+-----+-----+
|      smurf.|280790|
|      neptune.|107201|
|      normal.| 97278|
|      back.| 2203|
|      satan.| 1589|
|      ipsweep.| 1247|
|      portsweep.| 1040|
|      warezclient.| 1020|
|      teardrop.| 979|
|      pod.| 264|
|      nmap.| 231|
|      guess_passwd.| 53|
|      buffer_overflow.| 30|
|      land.| 21|
|      warezmaster.| 20|
|      imap.| 12|
|      rootkit.| 10|
|      loadmodule.| 9|
|      ftp_write.| 8|
|      multihop.| 7|
+-----+-----+

```

only showing top 20 rows

Another way is to write the code yourself to do it. You can extract the aggregated data as a `pandas` `DataFrame` and then plot it as a regular bar chart.

```
1 labels_df = pd.DataFrame(labels.toPandas())
2 labels_df.set_index("label", drop=True, inplace=True)
3 labels_fig = labels_df.plot(kind='barh')
4
5 plt.rcParams["figure.figsize"] = (7, 5)
6 plt.rcParams.update({'font.size': 10})
7
8
```



Connections based on protocols and attacks

Let's look at which protocols are most vulnerable to attacks now based on the following SQL query.

```
1 attack_protocol = sqlContext.sql("""
2     SELECT
3         protocol_type,
4         CASE label
5             WHEN 'normal.' THEN 'no attack'
6             ELSE 'attack'
7         END AS state,
8         COUNT(*) as freq
9     FROM connections
10    GROUP BY protocol_type, state
11    ORDER BY 3 DESC
12    """)
13 attack_protocol.show()
```



```
+-----+-----+-----+
|protocol_type|  state|  freq|
+-----+-----+-----+
|      icmp|  attack|282314|
|      tcp|  attack|113252|
|      tcp|no attack| 76813|
|      udp|no attack| 19177|
|      icmp|no attack|  1288|
|      udp|  attack|   1177|
+-----+-----+-----+
```

Well, looks like ICMP connections followed by TCP connections have had the maximum attacks!

Connection stats based on protocols and attacks

Let's take a look at some statistical measures pertaining to these protocols and attacks for our connection requests.

```
1 attack_stats = sqlContext.sql("""
2     SELECT
```

```

3         protocol_type,
4         CASE label
5             WHEN 'normal.' THEN 'no attack'
6             ELSE 'attack'
7         END AS state,
8         COUNT(*) as total_freq,
9         ROUND(AVG(src_bytes), 2) as mean_src_bytes,
10        ROUND(AVG(dst_bytes), 2) as mean_dst_bytes,
11        ROUND(AVG(duration), 2) as mean_duration,
12        SUM(num_failed_logins) as total_failed_logins,
13        SUM(num_compromised) as total_compromised,
14        SUM(num_file_creations) as total_file_creations,
15        SUM(su_attempted) as total_root_attempts,
16        SUM(num_root) as total_root_acceses
17    FROM connections
18    GROUP BY protocol_type, state
19    ORDER BY 3 DESC
20    "")
21 attack_stats.show()

```

protocol_type	state	total_freq	mean_src_bytes	mean_dst_bytes	mean_duration	total_failed_logins	total_compromised	total_file_cre
icmp	attack	282314	932.14	0.0	0.0	0	0	
tcp	attack	113252	9880.38	881.41	23.19	57	2269	
tcp	no attack	76813	1439.31	4263.97	11.08	18	2776	
udp	no attack	19177	98.01	89.89	1054.63	0	0	
icmp	no attack	1288	91.47	0.0	0.0	0	0	
udp	attack	1177	27.5	0.23	0.0	0	0	

Looks like average amount of data being transmitted in TCP requests are much higher which is not surprising. Interestingly attacks have a much higher average payload of data being transmitted from the source to the destination.

✓ Filtering connection stats based on the TCP protocol by service and attack type

Let's take a closer look at TCP attacks given that we have more relevant data and statistics for the same. We will now aggregate different types of TCP attacks based on service, attack type and observe different metrics.

```

1 tcp_attack_stats = sqlContext.sql("""
2     SELECT
3         service,
4         label as attack_type,
5         COUNT(*) as total_freq,
6         ROUND(AVG(duration), 2) as mean_duration,
7         SUM(num_failed_logins) as total_failed_logins,
8         SUM(num_file_creations) as total_file_creations,
9         SUM(su_attempted) as total_root_attempts,
10        SUM(num_root) as total_root_acceses
11    FROM connections
12    WHERE protocol_type = 'tcp'
13    AND label != 'normal.'
14    GROUP BY service, attack_type
15    ORDER BY total_freq DESC
16    """)
17 tcp_attack_stats.show()

```

service	attack_type	total_freq	mean_duration	total_failed_logins	total_file_creations	total_root_attempts	total_root_acceses
private	neptune.	101317	0.0	0	0	0.0	0
http	back.	2203	0.13	0	0	0.0	0
other	satan.	1221	0.0	0	0	0.0	0
private	portsweep.	725	1915.81	0	0	0.0	0
ftp_data	warezclient.	708	403.71	0	0	0.0	0
ftp	warezclient.	307	1063.79	0	0	0.0	0
other	portsweep.	260	1058.22	0	0	0.0	0
telnet	neptune.	197	0.0	0	0	0.0	0
http	neptune.	192	0.0	0	0	0.0	0
finger	neptune.	177	0.0	0	0	0.0	0
ftp_data	neptune.	170	0.0	0	0	0.0	0
private	satan.	170	0.05	0	0	0.0	0
csnet_ns	neptune.	123	0.0	0	0	0.0	0
smtp	neptune.	120	0.0	0	0	0.0	0
remote_job	neptune.	118	0.0	0	0	0.0	0
pop_3	neptune.	118	0.0	0	0	0.0	0
discard	neptune.	115	0.0	0	0	0.0	0
iso_tsap	neptune.	115	0.0	0	0	0.0	0

	systat	neptune.	113	0.0	0	0	0.0	0
	domain	neptune.	112	0.0	0	0	0.0	0

only showing top 20 rows

There are a lot of attack types and the preceding output shows a specific section of the same.

Filtering connection stats based on the TCP protocol by service and attack type

We will now filter some of these attack types by imposing some constraints based on duration, file creations, root accesses in our query.

```
1 tcp_attack_stats = sqlContext.sql("""
2     SELECT
3         service,
4         label as attack_type,
5         COUNT(*) as total_freq,
6         ROUND(AVG(duration), 2) as mean_duration,
7         SUM(num_failed_logins) as total_failed_logins,
8         SUM(num_file_creations) as total_file_creations,
9         SUM(su_attempted) as total_root_attempts,
10        SUM(num_root) as total_root_accesses
11    FROM connections
12    WHERE (protocol_type = 'tcp'
13           AND label != 'normal.')
14    GROUP BY service, attack_type
15    HAVING (mean_duration >= 50
16           AND total_file_creations >= 5
17           AND total_root_accesses >= 1)
18    ORDER BY total_freq DESC
19    """)
20 tcp_attack_stats.show()
```

service	attack_type	total_freq	mean_duration	total_failed_logins	total_file_creations	total_root_attempts	total_root_accesses
telnet	buffer_overflow.	21	130.67	0	15	0.0	5
telnet	loadmodule.	5	63.8	0	9	0.0	3
telnet	multihop.	2	458.0	0	8	0.0	93

Interesting to see multihop attacks being able to get root accesses to the destination hosts!

Subqueries to filter TCP attack types based on service

Let's try to get all the TCP attacks based on service and attack type such that the overall mean duration of these attacks is greater than zero (> 0). For this we can do an inner query with all aggregation statistics and then extract the relevant queries and apply a mean duration filter in the outer query as shown below.